

Tutorial 4

Statistical Computation and Analysis
Spring 2025

Tutorial Outline

- Probabilistic Programming Languages
- Analyzing the posterior
 - Arviz
 - KDE
 - HDI
 - Savage-Dickey Density Ratio
 - ROPE
- PyTensor


Probabilistic Programming Languages

- Challenge
 - Fully probabilistic models often lead to analytically intractable expressions
- Probabilistic Programming Languages
 - Allow clear separation between model creation and inference
 - Users specify a full probabilistic model by writing a few lines of code, and then inference follows automatically

Translating a Graphical Model to a PPL

- Coin flipping

Graphical model



Equations

$$\theta \sim \text{Beta}(\alpha, \beta)$$
$$y \sim \text{Binom}(\theta, N)$$

PyMC (a PPL)

```
with pm.Model() as our_first_model:  
     $\theta$  = pm.Beta('theta', alpha=1., beta=1.)  
    y = pm.Bernoulli('y', p= $\theta$ , observed=data)  
    idata = pm.sample(1000, random_seed=4591)
```

Translating a Graphical Model to a PPL

■ Coin flipping

```
with pm.Model() as our_first_model:  
     $\theta$  = pm.Beta('θ', alpha=1., beta=1.)  
    y = pm.Bernoulli('y', p= $\theta$ , observed=data)  
    idata = pm.sample(1000, random_seed=4591)
```

Line 1: creates a container for our model. Everything inside the with block will be automatically added to our_first_model.

Line 2: prior

Line 3: likelihood


Line 4: computes the posterior using numerical methods.

Translating a Graphical Model to a PPL

- The use of numerical methods solves the challenge of posteriors we cannot compute.
- However, what we get are samples from the posterior, rather than the posterior itself.
- Every time we run the inference, the samples will change.
- **If the inference process works as expected, the samples will be representative of the posterior distribution, and we will obtain the same conclusion from any of those samples.**

Translating a Graphical Model to a PPL


- Number of emails
 - Poisson likelihood
 - Gamma prior on λ



Translating a Graphical Model to a PPL

- Number of emails

Graphical model



Equations

$$\lambda \sim \text{Gamma}(\alpha, \beta)$$
$$k \sim \text{Poisson}(\lambda)$$

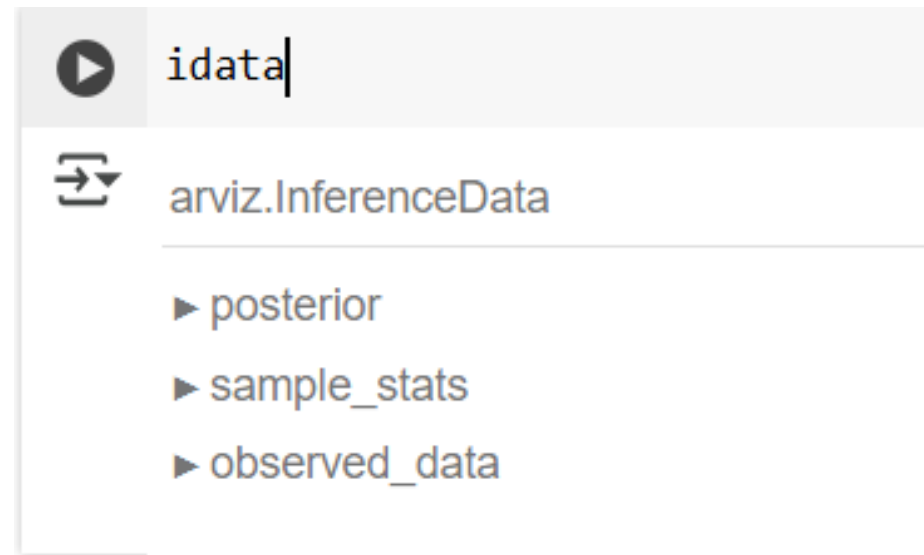
PyMC (a PPL)

```
coords = {"data": np.arange(n)}

with pm.Model(coords = coords) as our_first_model:
    lambda_ = pm.Gamma('lam', alpha = 1.68, beta = 0.0569)
    k = pm.Poisson('k', mu = lambda_, observed=data, dims = 'data')
    idata = pm.sample(1000, chains = 4)
```


Inference Data Object

- idata – what we get back from the inference process
- Container for all the data generated by PyMC.



- Xarray object
- We will use the Arviz library to analyze the posterior

Inference Data Object

- Observed data

- posterior
- sample_stats
- observed_data

xarray.Dataset

► Dimensions: (data: 24)

Number of samples collected (24)

▼ Coordinates:

data (data) int64 0 1 2 3 4 5 6 ... 18 19 20 21 22 23

Their coordinates – here just 0-23

▼ Data variables:

k (data) int64 10 5 2 11 9 7 8 ... 9 6 12 8 6 8 6


The values – the numbers of emails received by each student


► Indexes: (1)

► Attributes: (4)

Inference Data Object

■ Posterior

 idata

 arviz.InferenceData

▼ posterior





xarray.Dataset

► Dimensions:



(chain: 4, draw: 1000)

4 chains of 1000 samples (draws) each

▼ Coordinates:

chain	(chain)	int64	0 1 2 3	 
draw	(draw)	int64	0 1 2 3 4 5 ... 995 996 997 998 999	 

▼ Data variables:

lambda	(chain, draw)	float64	9.073 9.011 7.855 ... 8.935 8.721	 
--------	---------------	---------	-----------------------------------	---



► Indexes: (2)

► Attributes: (6)

The lambda values (4, 1000)


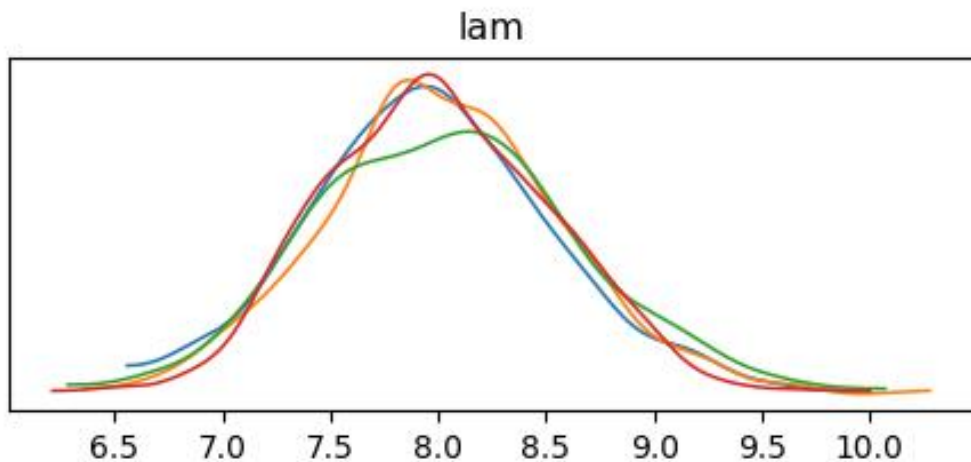
Analyzing the Posterior

- Look at the results
 - `az.plot_trace(idata)`





Analyzing the Posterior

- Left: the four chains showing the values of the posterior distribution (**KDE plot = Kernel Density Estimation**)
- Right: For now – the overlap between the four colors shows that the sampler is working.




Analyzing the Posterior

- Compare our analytical results (from Tutorial 3) to these:



Analyzing the Posterior

- We want all the chains to overlap as much as possible.
 - How much?
 - We'll learn statistics with thresholds we can calculate
 - For now, It looks pretty good.




Analyzing the Posterior

- How close are the means of the different chains?


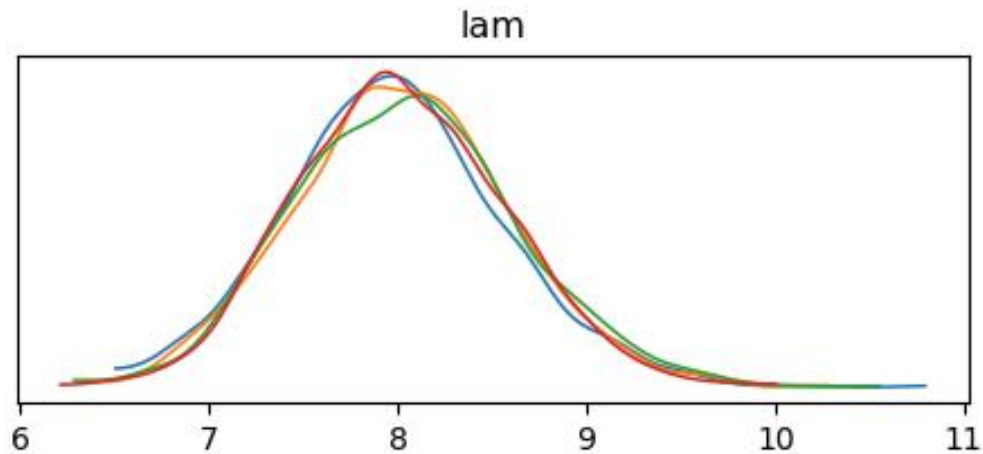
```
ms = idata.posterior.mean(dim = 'draw')  
print(f'The means of the 4 chains are: {np.round(ms.lam.to_numpy(), 3)}')
```

→ The means of the 4 chains are: [7.97 8.028 8.036 7.998]



Analyzing the Posterior

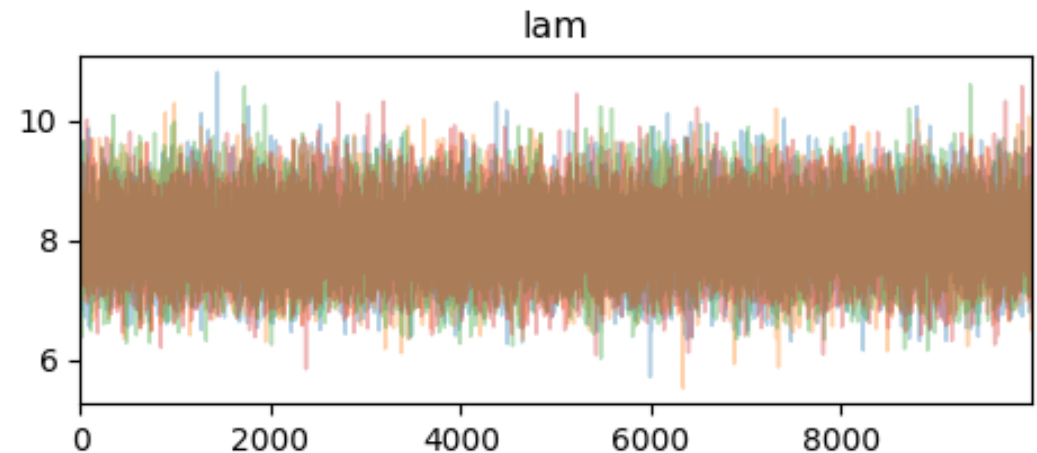

- How can we get the chains to look more similar?
 - Sample more (2000)



The means of the 4 chains are: [7.984 8.029 8.043 8.016]

Analyzing the Posterior

- How can we get the chains to look more similar?
 - Sample more (10,000)



The means of the 4 chains are: [8.002 8.013 8.018 8.016]



Highest Density Interval (HDI)

- We computed the mean
- We can also compute other statistics, such as the HDI
- The Z% HDI is the interval of values containing Z% of the probability

$$\int_{x \in Z\% \text{ HDI}} dx \, p(x) = \frac{Z}{100}$$

Highest Density Interval (HDI)


- There are a lot of intervals like this.
- The HDI is the shortest interval containing $Z\%$ of the distribution.
- Equivalent: the probability of all values inside the HDI are higher than those out of it.



Highest Density Interval (HDI)

- The default in Arviz is the 94% HDI, so that's what we'll generally use.
- Often, the 50% HDI is also reported
- Before we plot, we can in general combine the 4 chains to 4000 samples:

The total mean of the 4000 samples is: 8.095



Highest Density Interval (HDI)







- We can also get the data as samples instead of as separate chains using `az.extract`

```
#we can also get all the data as 4000 samples using extract, instead of as four separate chains
az.extract(idata)
|
#if you specify number of samples, you get random sampling out of the 4000
az.extract(idata, num_samples = 3500)
```

xarray.Dataset

► Dimensions: (sample: 3500)

▼ Coordinates:

sample	(sample)	object	MultiIndex	 
chain	(sample)	int64	2 1 2 0 1 1 2 0 ... 2 1 1 3 3 2 0 1	 
draw	(sample)	int64	961 68 332 587 ... 208 289 484 635	 


▼ Data variables:

lam	(sample)	float64	9.237 8.037 7.363 ... 8.934 7.618	 
-----	----------	---------	-----------------------------------	---

Highest Density Interval (HDI)

- We can plot HDI and compute the values

```
#plotting the posterior with the HDI  
az.plot_posterior(idata)
```



```
#and getting the values in a table  
az.summary(idata, kind = 'stats').round(2)
```

	mean	sd	hdi_3%	hdi_97%
lam	8.1	0.57	7.12	9.27



Lower
limit

Upper
limit

Help Pages


- Arviz: <https://python.arviz.org/en/stable/api/index.html>
- PyMC: <https://www.pymc.io/projects/docs/en/stable/api.html>
- Preliz: <https://preliz.readthedocs.io/en/latest/index.html>

Posterior Based Decisions

- Sometimes we would like to do more than describe the posterior
- Sometimes we'd like to make a decision
 - Is the coin fair?
 - Is the person sick?
- Do biomedical students receive 9 emails a day?

Posterior Based Decisions


- 9 is inside the 94% HDI.



- We can compute a value called the **Savage-Dickey density ratio**.

Savage-Dickey Density Ratio

- The ratio of the posterior and prior densities at that value.
 - Evaluates how much support the posterior provides for a given value
- $BF_{01} = 7.97$
 - The value of 9 emails is 7.97 times more likely under the posterior than under the prior.
- $BF_{10} = 1/BF_{01}$



Savage-Dickey Density Ratio

- How do we a decision?
 - Rule of thumb

BF_{01}	Interpretation
< 3.2	Not worth mentioning
3.2 to 10	Substantial
10 to 100	Strong
> 100	Decisive

ROPE

- Region of Practical Equivalence
 - For example, students will not each get exactly 9 emails a day.
 - Let's say we think between 8-10 is pretty much the same as 9.
 - ROPE is defined using domain knowledge.

ROPE


- If the ROPE and HDI do not overlap
 - Students don't receive 9 emails a day
- If the ROPE contains the entire HDI
 - Students receive 9 emails a day
- If the ROPE and HDI partially overlap
 - We can't say if students do or do not receive 9 emails a day

ROPE

- If the ROPE and HDI do not overlap
 - Students don't receive 9 emails a day
- If the ROPE contains the entire HDI
 - Students receive 9 emails a day
- If the ROPE and HDI partially overlap
 - We can't say if students do or do not receive 9 emails a day
- If the definition of the ROPE changes our conclusions – maybe question it.


ROPE

- Can we conclude if students receive 9 emails a day?
 - No




ROPE


(A) Decision: Reject Null Value




(D) Decision: Undecided




(B) Decision: Accept Null Value




(E) Decision: Undecided



(C) Decision: Accept Null Value



(F) Decision: Undecided



PyTensor

- Allows for defining, optimizing, and efficiently evaluating mathematical expressions involving multi-dimensional arrays.
- A tensor is a generalized mathematical structure that can represent scalars, vectors, matrices, and higher-dimensional arrays.
- Unlike other variables, tensors don't have to have a defined shape to be manipulated.
- This can be useful for cases in which tensors exist without a concrete shape until execution or have dimensions that vary per sample.
 - Used in deep learning.
- Computational graphs enable automatic differentiation for efficient gradient computation using the chain rule.