

Tutorial 6

Statistical Computation and Analysis
Spring 2025

Tutorial Outline

- Grid Method
- MCMC
 - Metropolis-Hastings
- Diagnostics
 - Convergence
 - Divergence

Grid Method

- $P(\text{model}|\text{data}) = \frac{P(\text{data}|\text{model})P(\text{model})}{\sum_{\text{models}} P(\text{data}|\text{model})P(\text{model})}$
- The marginal likelihood is often a challenging and computationally expensive integral to calculate.
- If we know that prior and the likelihood at a certain point
 - We can multiply them and get something proportional to the posterior
- If we know the prior and likelihood at a bunch of points
 - We can multiply them, and then normalize, and get an estimation of the posterior


Grid Method

- In the past we looked at the question: How many emails do you receive each day?
- Likelihood:
- For n iid observations:
 - $P(data|\lambda) = \prod_{i=1}^n \frac{\lambda^{k_i} e^{-\lambda}}{k_i!} = \frac{\lambda^{\sum_{i=1}^n k_i} e^{-n\lambda}}{\prod_{i=1}^n k_i!}$
 - $S = \sum_{i=1}^n k_i$ (total number of observed events)
 - $P(data|\lambda) = \prod_{i=1}^n \frac{\lambda^{k_i} e^{-\lambda}}{k_i!} = \frac{\lambda^S e^{-n\lambda}}{\prod_{i=1}^n k_i!}$

Grid Method

- Prior:

- $P(\lambda) \sim \text{Gamma}(\alpha, \beta) = \frac{\beta^\alpha \lambda^{\alpha-1} e^{-\beta\lambda}}{\Gamma(\alpha)}$



Grid Method

- Posterior:
- $P(\lambda|\text{data}) \propto P(\text{data}|\lambda)P(\lambda)$
- $P(\lambda|\text{data}) \sim \text{Gamma}(\alpha + S, n + \beta)$
 - S = total number of emails received
 - n = Number of students
- We showed this analytically.
- Let's compare this to the results using the grid method.


Grid Method

Steps:

1. Define a reasonable interval for the parameter (the prior should give you a hint).
 - Let's place 100 points between 0-40.
2. Place a grid of points (generally equidistant) on that interval.
3. For each point in the grid, multiply the likelihood and the prior.

Grid Method

- We can see the overlap between the grid method and the analytical results.
- We wasted a lot of our grid sampling values below 5 and above 10.



Grid Method

- Disadvantages:
 - We focus equally on all areas of the grid, and not on high probability areas.
 - Curse of dimensionality




Figure 10.2: A grid with the same resolution in 1, 2, and 3 dimensions

Markovian Chain Monte Carlo Methods (MCMC)

- Allow us to get samples from the posterior distribution as long as we can compute the likelihood and the prior pointwise.
- Same condition we needed for the grid method.
- MCMC methods can efficiently sample from higher-probability regions in very high dimensions.
- MCMC methods visit each region of the parameter space following their relative probabilities.
 - If the probability of region A is twice that of region B, we will obtain twice as many samples from A as we will from B.
- Even if we are not capable of computing the whole posterior analytically, we could use MCMC methods to take samples from it.
- MCMC will give us samples from the correct distribution asymptotically (for an infinite number of samples).

Markovian Chain Monte Carlo Methods (MCMC)

- **Monte Carlo methods:**

- A broad family of algorithms that use random sampling to compute or simulate a given process.

- **Markovian methods:**

- Techniques or models based on the Markov property, which is the idea that the future state of a system depends only on its current state, and not on its previous history.
- Future outcomes are independent of past events given the present state.

Markovian Chain Monte Carlo Methods (MCMC)

- **Markov Chain:**

- A mathematical object that consists of a sequence of states and a set of transition probabilities that describe how to move among the states.
- Simple example: flip a coin and if you get heads take a step to the right, otherwise step to the left.
- A chain is Markovian if the probability of moving to any other state depends only on the current state.

Metropolis-Hastings

- Enables us to obtain samples from any probability distribution given that we can compute at least a value proportional to it.
 - Thus ignoring the normalization factor.
 - Useful since often the harder part is to compute the normalization factor.

Metropolis-Hastings

- We want to estimate the depth of the bottom of the lake, and find the deepest point.
 - Lake is muddy so we can't see.
- 1. Take a boat and a stick.
- 2. Choose a random place in the lake and move the boat there.
- 3. Use the stick to measure the depth of the lake.
- 4. Move the boat to another point and take a new measurement.

Metropolis-Hastings

5. Compare the two measurements:

- If the new spot is deeper than the old one -> write it down.
- If the old spot was deeper -> either accept or reject the new spot.
 - Accept = write down the new spot.
 - Reject = go back to the old spot and write it down again.

6. Move to another point and repeat again and again.

Metropolis-Hastings

- What does this give us?
 - An estimate of the curvature of the bottom of the lake.
 - An analogy for the posterior.
 - The deepest point.
 - An analogy for the mode of the posterior.

Metropolis-Hastings

1. Choose an initial value for our parameter, x_i .
2. Write it down.
3. Suggest a new value for our parameter, x_{i+1} .

How do we choose x_{i+1} ?

- We use a **proposal distribution q** , such as a normal or uniform distribution.
- We sample from $q(x_{i+1}|x_i)$

Metropolis-Hastings

4. Compute the probability of accepting a new parameter value by using the **Metropolis-Hastings criteria**:

$$p_a(x_{i+1}|x_i) = \min\left(1, \frac{p(x_{i+1}) q(x_i|x_{i+1})}{p(x_i) q(x_{i+1}|x_i)}\right)$$

5. If $p_a(x_{i+1}|x_i) >$ random number from a continuous uniform distribution ranging $[0, 1]$: Accept the new parameter.
6. Otherwise – reject and stay where we are.
7. Repeat again and again until we have enough samples.

Metropolis-Hastings

Metropolis-Hastings criteria: $p_a(x_{i+1}|x_i) = \min(1, \frac{p(x_{i+1}) q(x_i|x_{i+1})}{p(x_i) q(x_{i+1}|x_i)})$

- The proposal distribution, q , is not part of the model. It is a component of the MCMC algorithm.
- If q is symmetric: $p_a(x_{i+1}|x_i) = \min(1, \frac{p(x_{i+1})}{p(x_i)})$

Metropolis-Hastings

- If q is symmetric: $p_a(x_{i+1}|x_i) = \min(1, \frac{p(x_{i+1})}{p(x_i)})$
- If x_{i+1} is a more probable state ($p(x_{i+1}) > p(x_i)$):
 - We for sure accept the new state.
 - Because then we get a probability of 1.
- If $p(x_{i+1}) < p(x_i)$:
 - We don't for sure move.
 - We move with a probability of $\frac{p(x_{i+1})}{p(x_i)} < 1$.

Metropolis-Hastings

- If $p(x_{i+1}) < p(x_i)$:
 - We move with a probability of $\frac{p(x_{i+1})}{p(x_i)} < 1$.

How do we quantify this?

- Choose a random number between 0-1 from a uniform distribution.
- If $\frac{p(x_{i+1})}{p(x_i)} = 0.8$
 - and our random number is 0.7 (there is a 70% chance of this)
 - Then we move, as we exceeded this value (the metropolis hastings criteria).

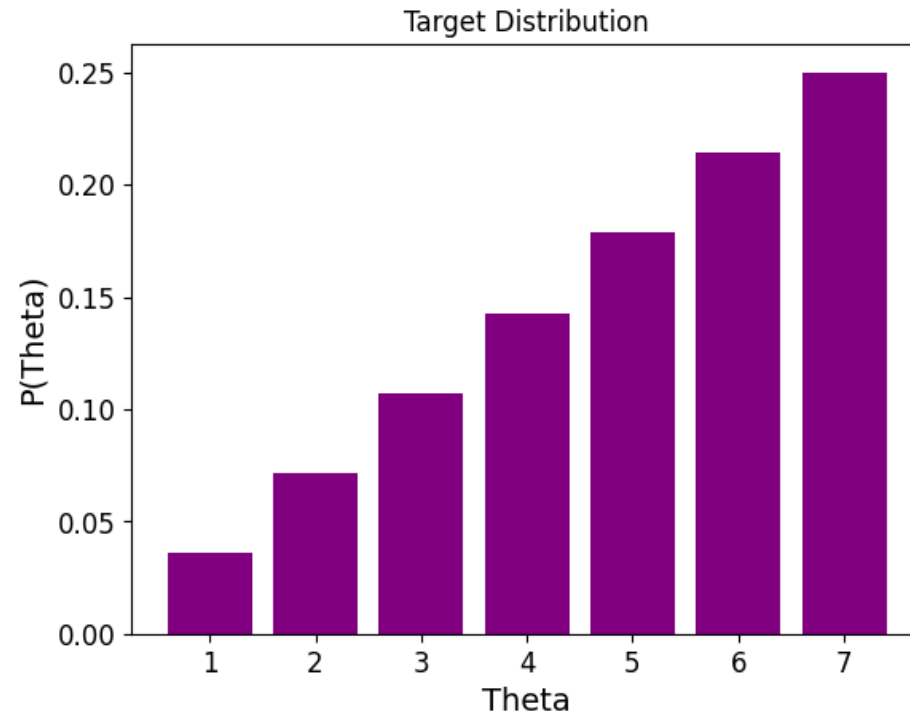
Metropolis-Hastings

- The island example.
 - We have 7 islands and we move between them.
 - The probability we are trying to sample from is $P(\theta) \propto \theta$
 - θ is the relative number of people on each island.




Metropolis-Hastings

- The island example.
 - The probability we are trying to sample from is $P(\theta) \propto \theta$
 - θ is the relative number of people on each island.




Metropolis-Hastings

- Start at a random island
 - We chose island 3
- So we write it down.




Metropolis-Hastings

- Let's try and move an island.
- We choose randomly and get island 4.
- $p_a(\theta_{i+1}|\theta_i) = \min\left(1, \frac{p(\theta_{i+1})}{p(\theta_i)}\right) = \min\left(1, \frac{4}{3}\right) = 1$

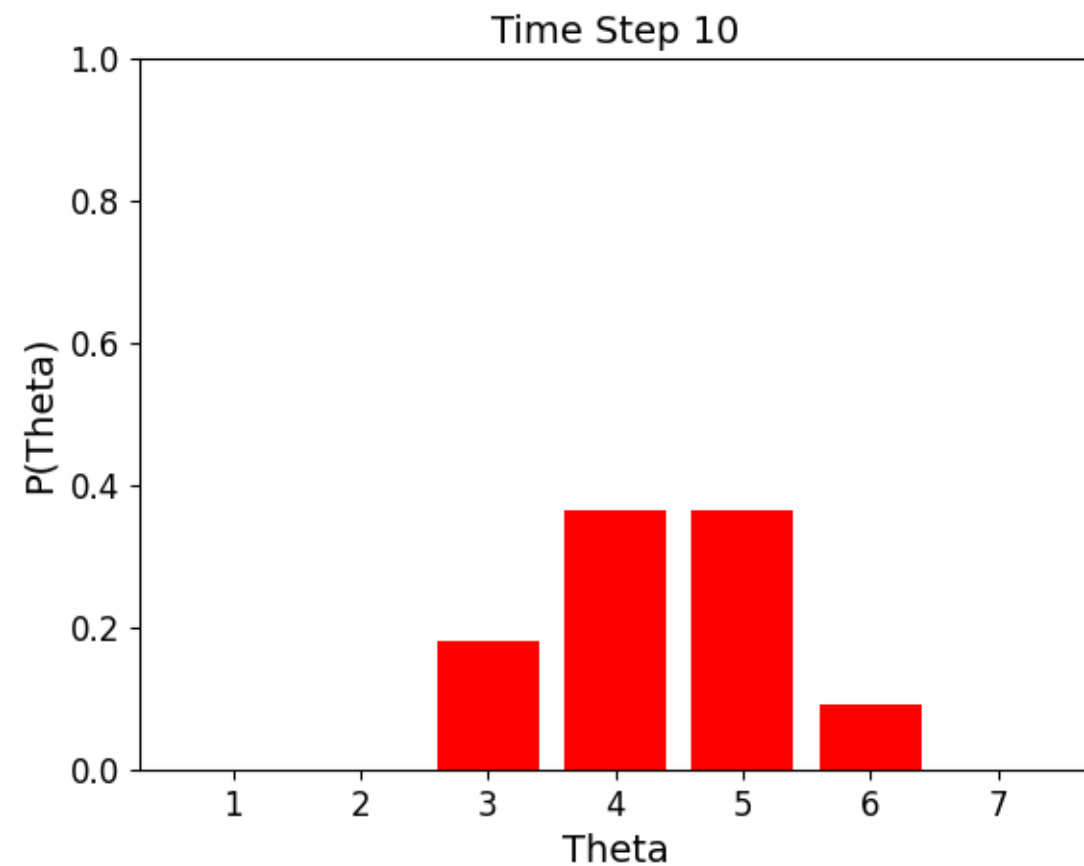
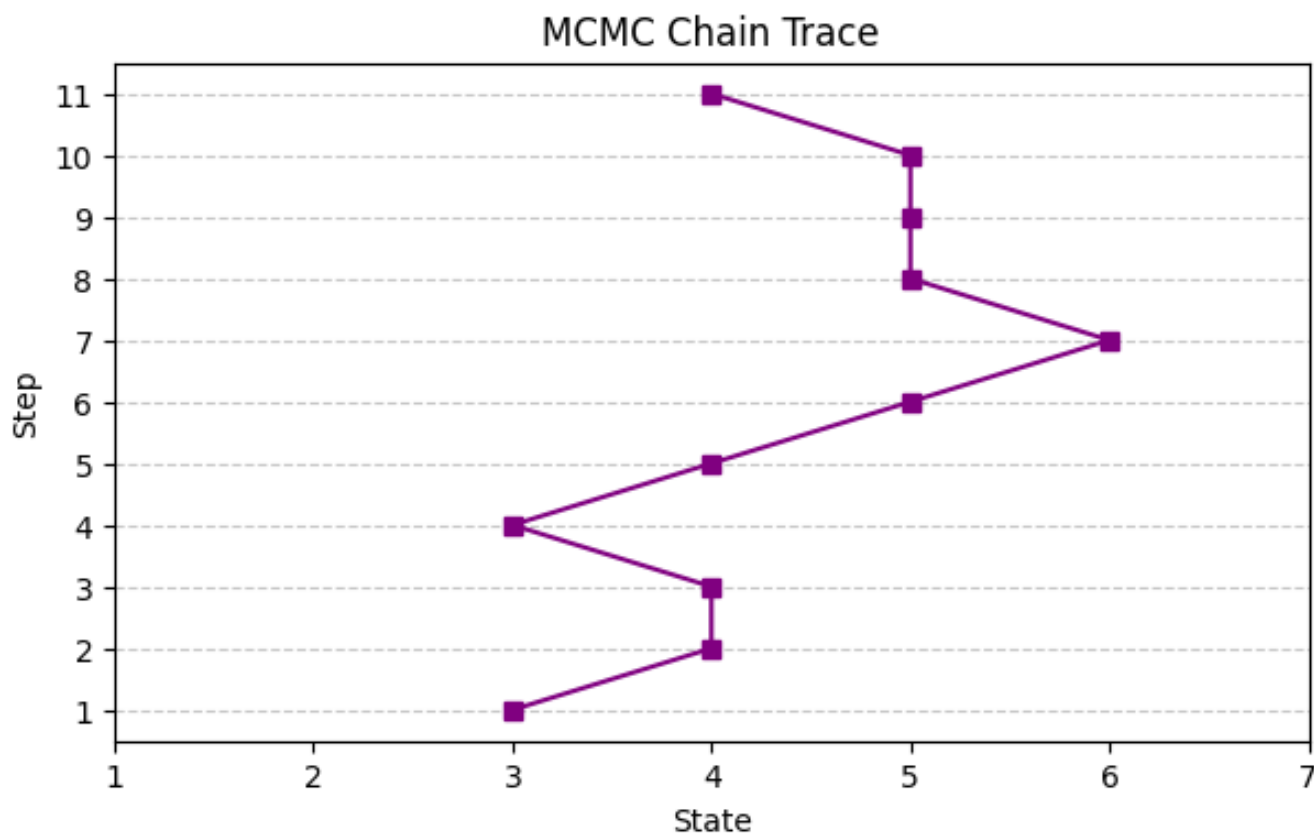


Metropolis-Hastings

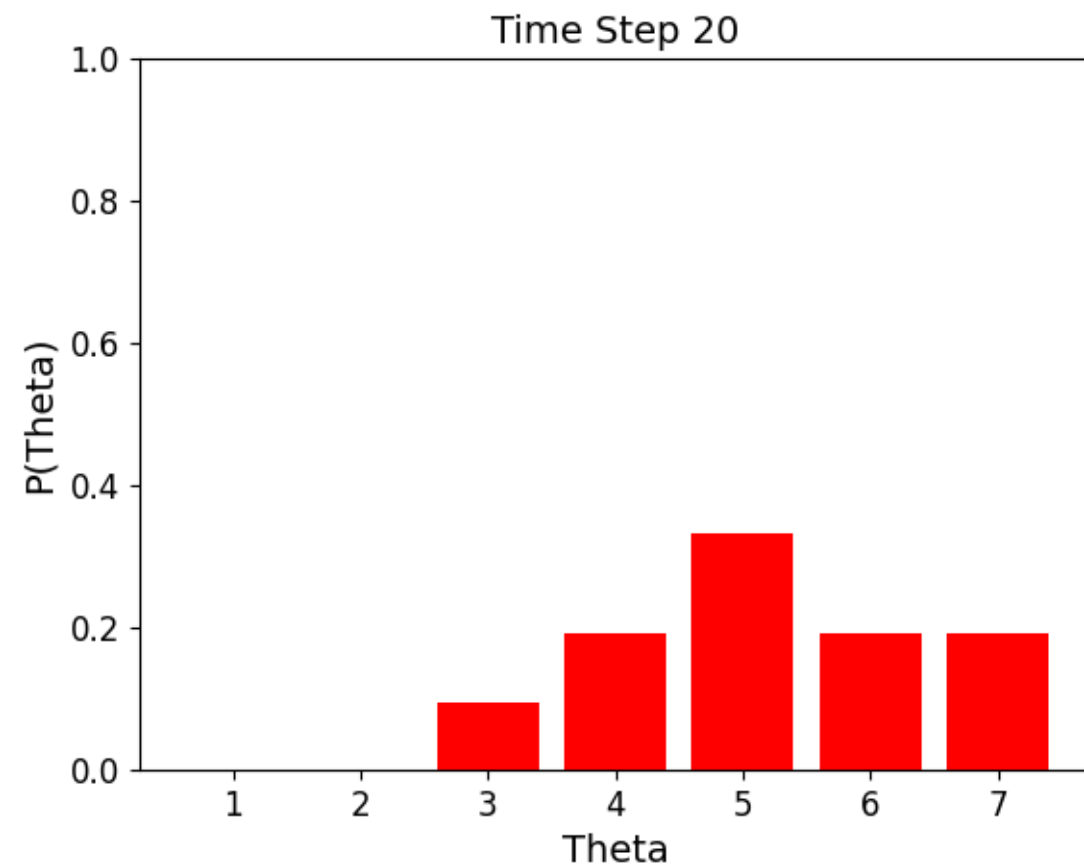
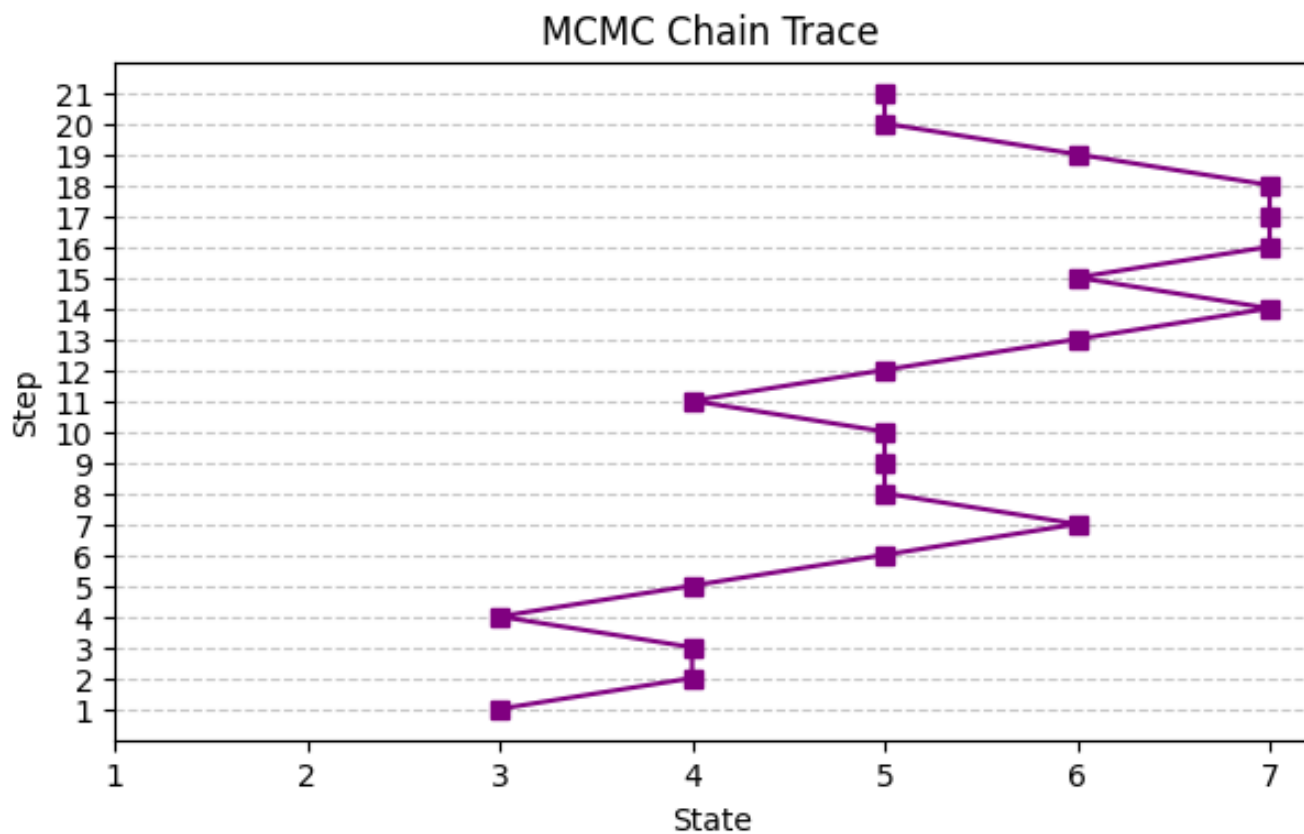
- Try to move back to 3
- $p_a(\theta_{i+1}|\theta_i) = \min\left(1, \frac{p(\theta_{i+1})}{p(\theta_i)}\right) \min\left(1, \frac{3}{4}\right) = \frac{3}{4}$
- Choose a random number (0.91)
- $0.75 < 0.91$ – don't move back to 3.
- Write 4 again.





Metropolis-Hastings



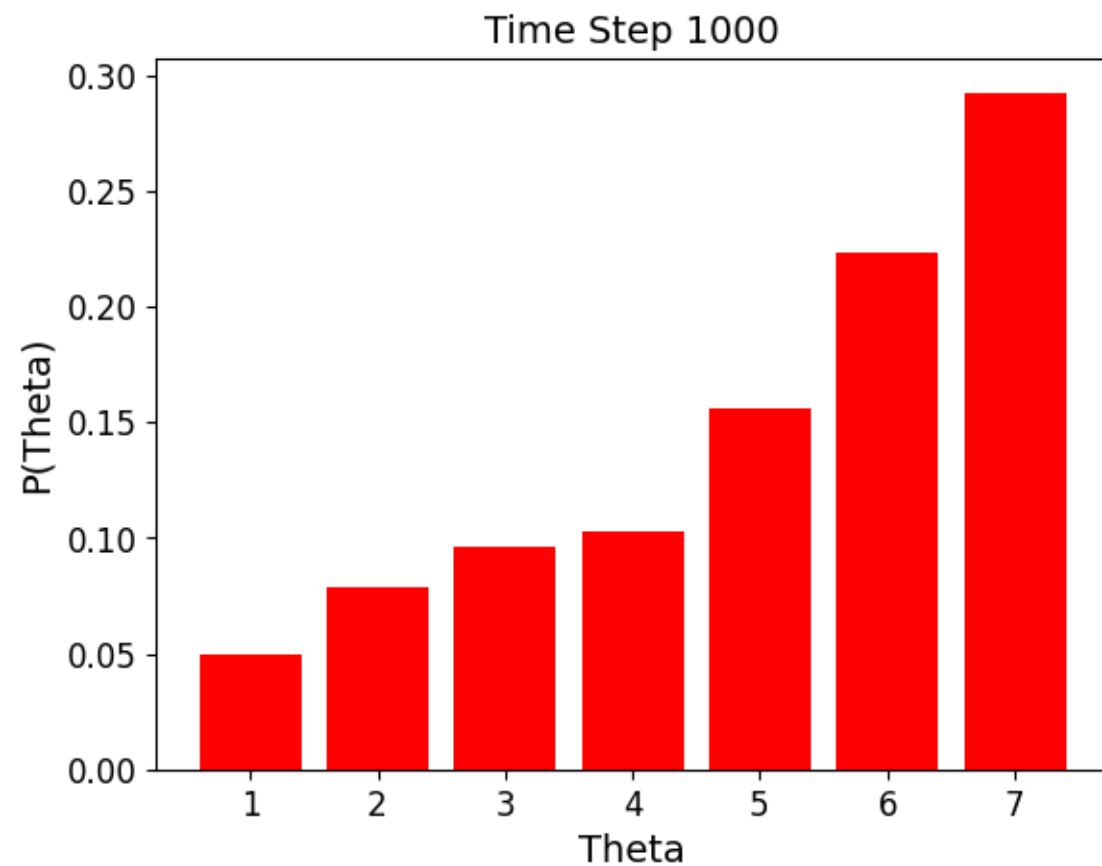
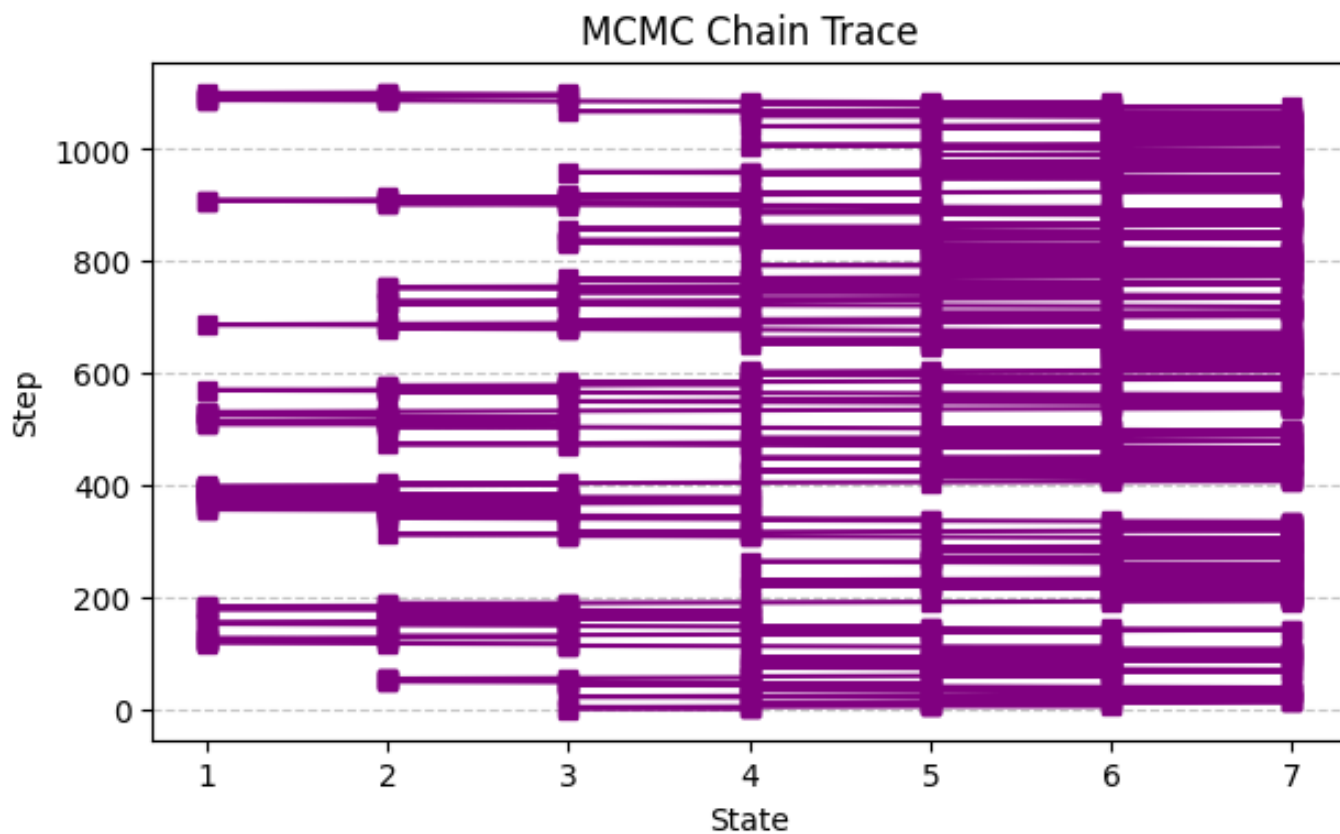
Metropolis-Hastings





Metropolis-Hastings



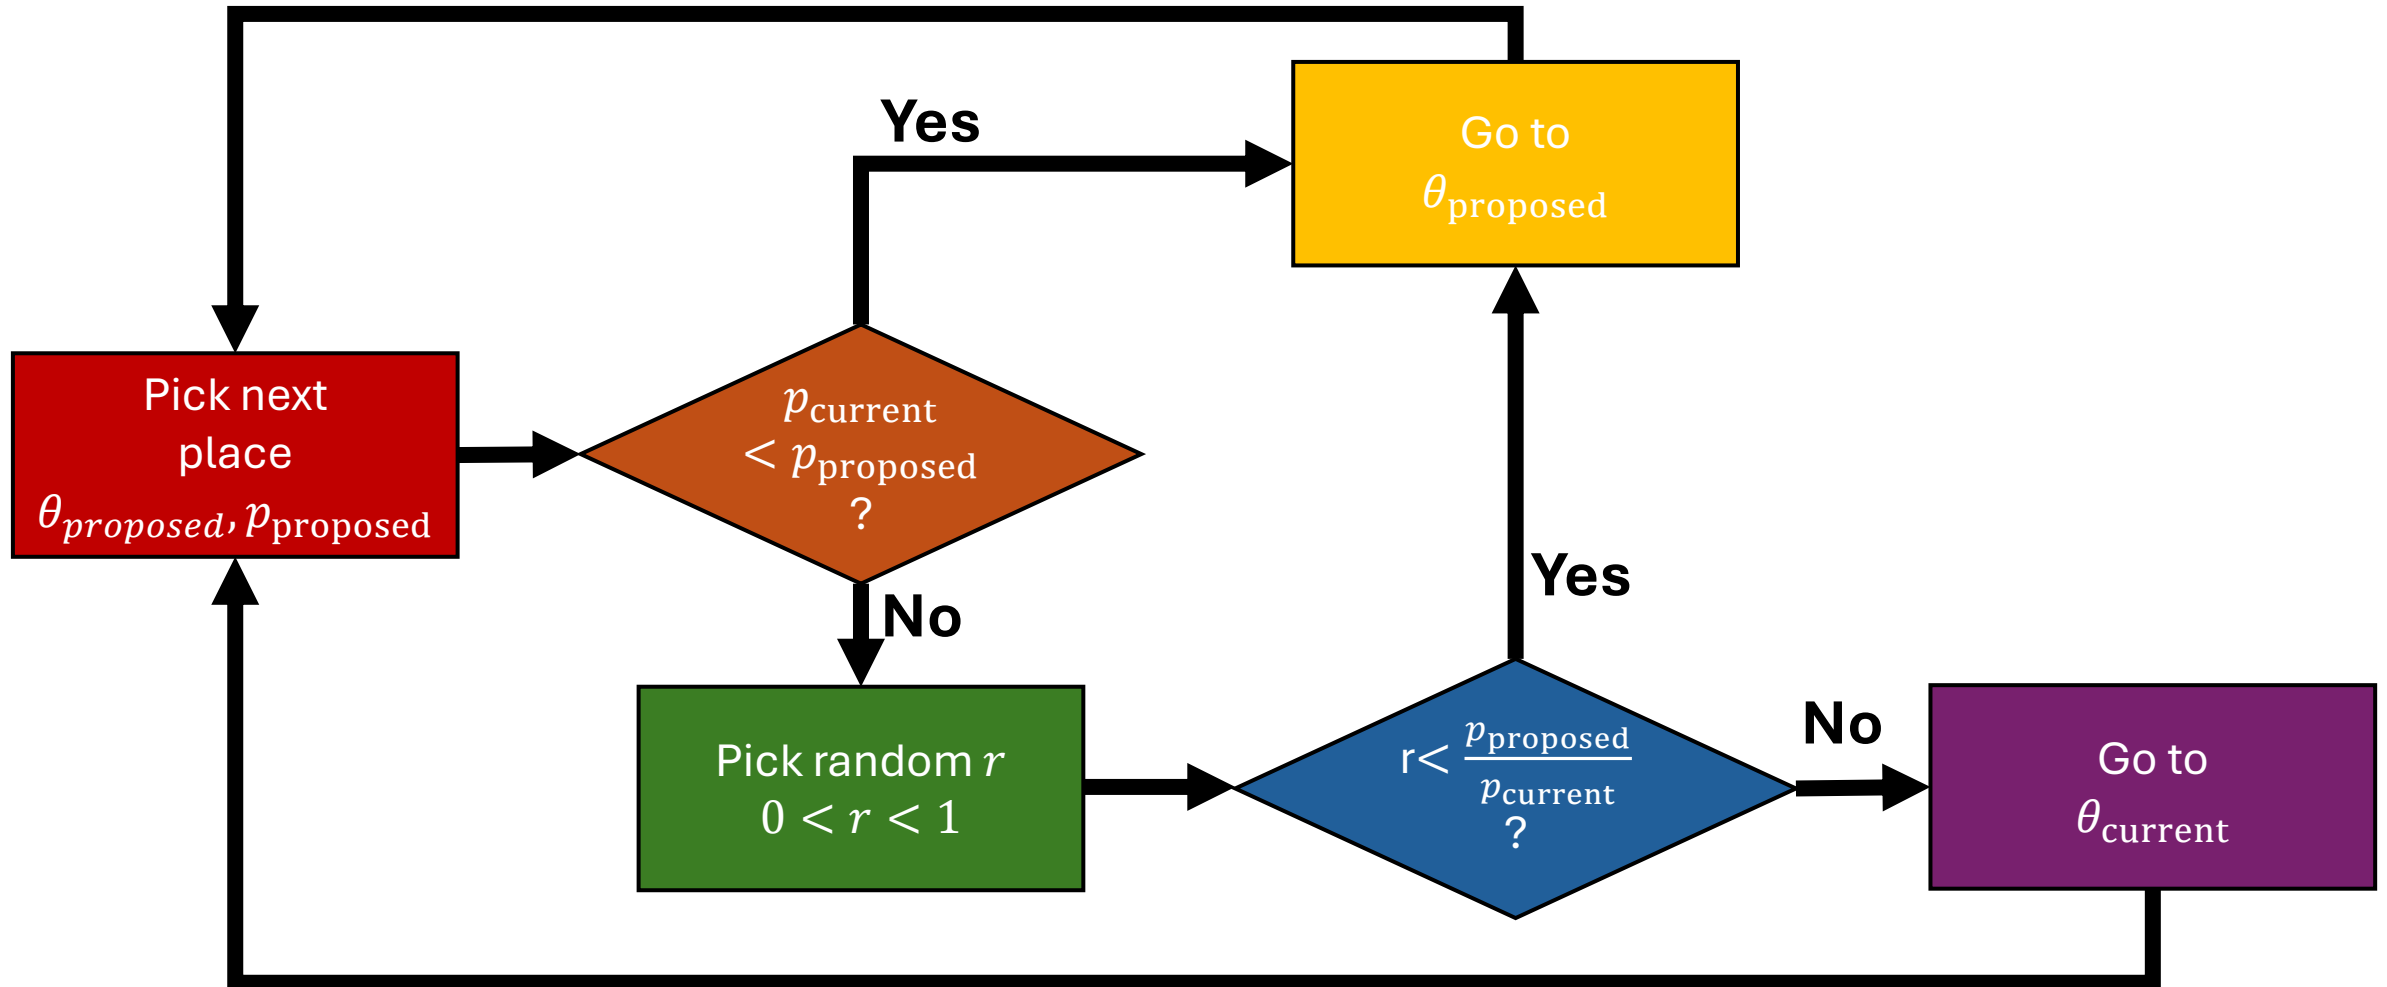
Metropolis-Hastings



Metropolis-Hastings



Metropolis-Hastings



Metropolis-Hastings

- How do we choose a proposal function?
 - The proposal function has to be scaled appropriately.
 - If we have a very broad proposal for a narrow posterior, all the proposals are to a place with no probability and we never accept a new state.
 - If your proposal is very narrow, get stuck near the starting point and explore that excessively.
 - We often use a normal distribution $N(0, \sigma)$ and tune the σ .
 - We can run the algorithm with an acceptance rate in mind (~70%-80%).
 - Look how many of the first N steps were accepted.
 - If the proposal is too narrow – we'll almost always accept.
 - We'll see we surpassed the acceptance rate and make the proposal wider.
 - If the proposal is too wide – we'll rarely accept.
 - We'll see we did not reach the acceptance rate and make the proposal narrower.

Diagnostics

- Using numerical methods means we are approximating the posterior with a finite number of samples.
 - Do we have a valid sample?
- We will use tests to see if the sampling algorithm worked.
 - These tests can spot problems with our samples.
 - They cannot prove that we have the correct distribution.
 - They can also provide evidence that the samples seem reasonable.

Three Different Priors For the Coin Flip

```
with pm.Model(coords=coords) as model_1:
    thet = pm.Uniform('thet', lower=0, upper=1)
    y = pm.Bernoulli('y', p=thet, observed=data, dims = 'data')
    idata1 = pm.sample(1000, chains = 4)

with pm.Model(coords=coords) as model_2:
    thet = pm.Uniform('thet', lower=-1, upper=2)
    y = pm.Bernoulli('y', p=thet, observed=data, dims = 'data')
    idata2 = pm.sample(1000, chains = 4)

with pm.Model(coords=coords) as model_3:
    thet = pm.Uniform('thet', lower=-0.1, upper=1)
    y = pm.Bernoulli('y', p=thet, observed=data, dims = 'data')
    idata3 = pm.sample(1000, chains = 4)
```

Three Different Priors For the Coin Flip

Progress	Draws	Divergences	Step size	Grad evals	Sampling Speed	Elapsed	Remaining
<div></div>	2000	0	1.46	1	2711.90 draws/s	0:00:00	0:00:00
<div></div>	2000	0	1.15	1	1374.80 draws/s	0:00:01	0:00:00
<div></div>	2000	0	1.16	3	923.66 draws/s	0:00:02	0:00:00
<div></div>	2000	0	1.16	3	688.67 draws/s	0:00:02	0:00:00

Progress	Draws	Divergences	Step size	Grad evals	Sampling Speed	Elapsed	Remaining
<div></div>	2000	92	1.20	1	2563.21 draws/s	0:00:00	0:00:00
<div></div>	2000	73	1.19	3	1302.72 draws/s	0:00:01	0:00:00
<div></div>	2000	92	1.05	3	865.75 draws/s	0:00:02	0:00:00
<div></div>	2000	86	1.76	1	644.52 draws/s	0:00:03	0:00:00

ERROR:pymc.stats.convergence:There were 343 divergences after tuning. Increase `target_accept` or reparameterize.

Progress	Draws	Divergences	Step size	Grad evals	Sampling Speed	Elapsed	Remaining
<div></div>	2000	13	1.23	3	2652.35 draws/s	0:00:00	0:00:00
<div></div>	2000	19	1.01	1	1383.77 draws/s	0:00:01	0:00:00
<div></div>	2000	9	1.85	1	833.34 draws/s	0:00:02	0:00:00
<div></div>	2000	7	1.73	1	574.27 draws/s	0:00:03	0:00:00

ERROR:pymc.stats.convergence:There were 48 divergences after tuning. Increase `target_accept` or reparameterize.



Convergence

- MCMC methods are guaranteed to converge for infinite samples.
- How do we know if we have a reasonable number of samples?
- The sampler has **converged** if the samples are stable.
 - Run the simulation multiple times and see if we get the same result every time.
 - We often used four chains.
 - This is also efficient as we can combine the samples from the different chains to have four times as many samples.


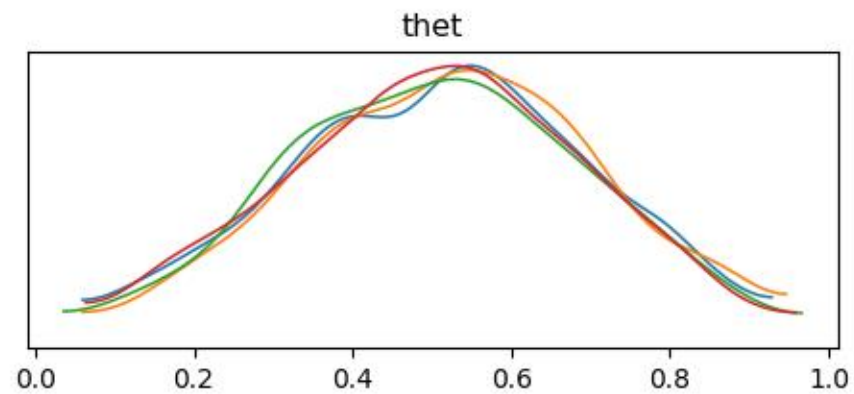
Convergence

- How do we check convergence?
 - Look at the trace plots and see if the chains look similar

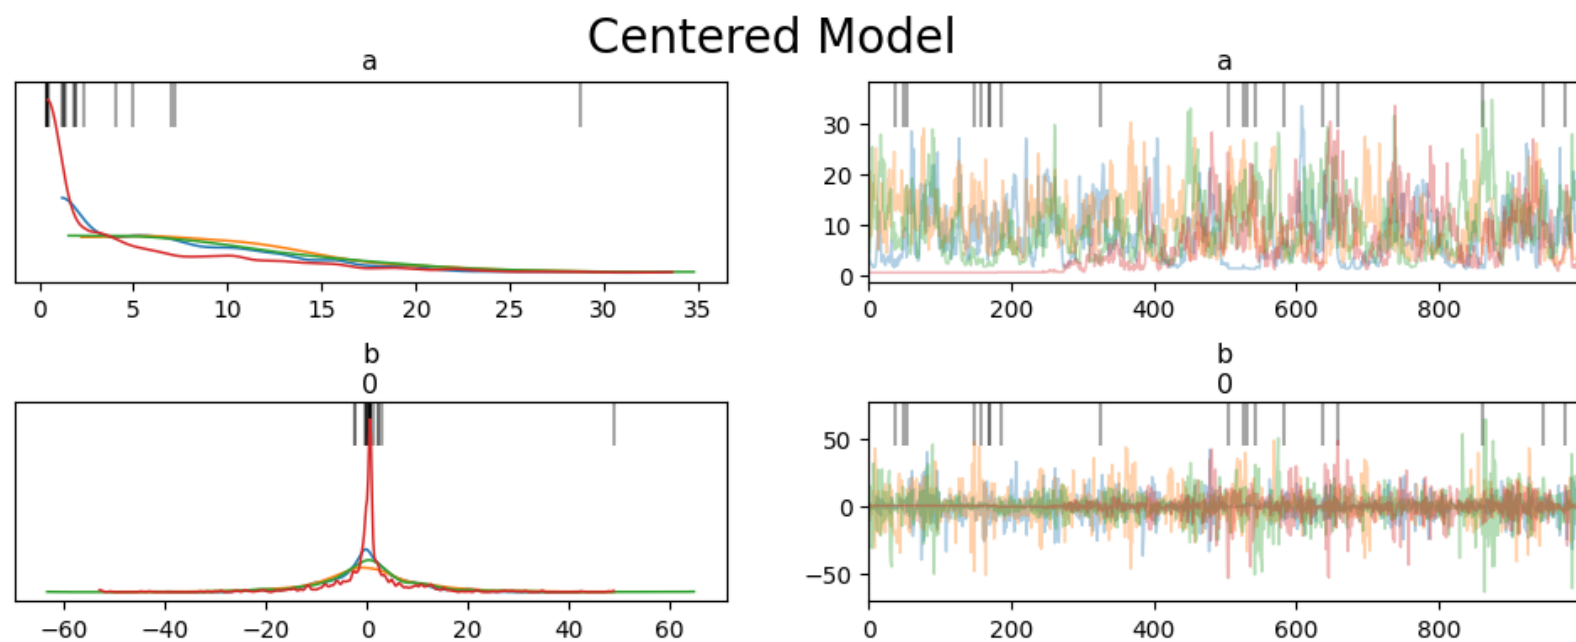
Less good:



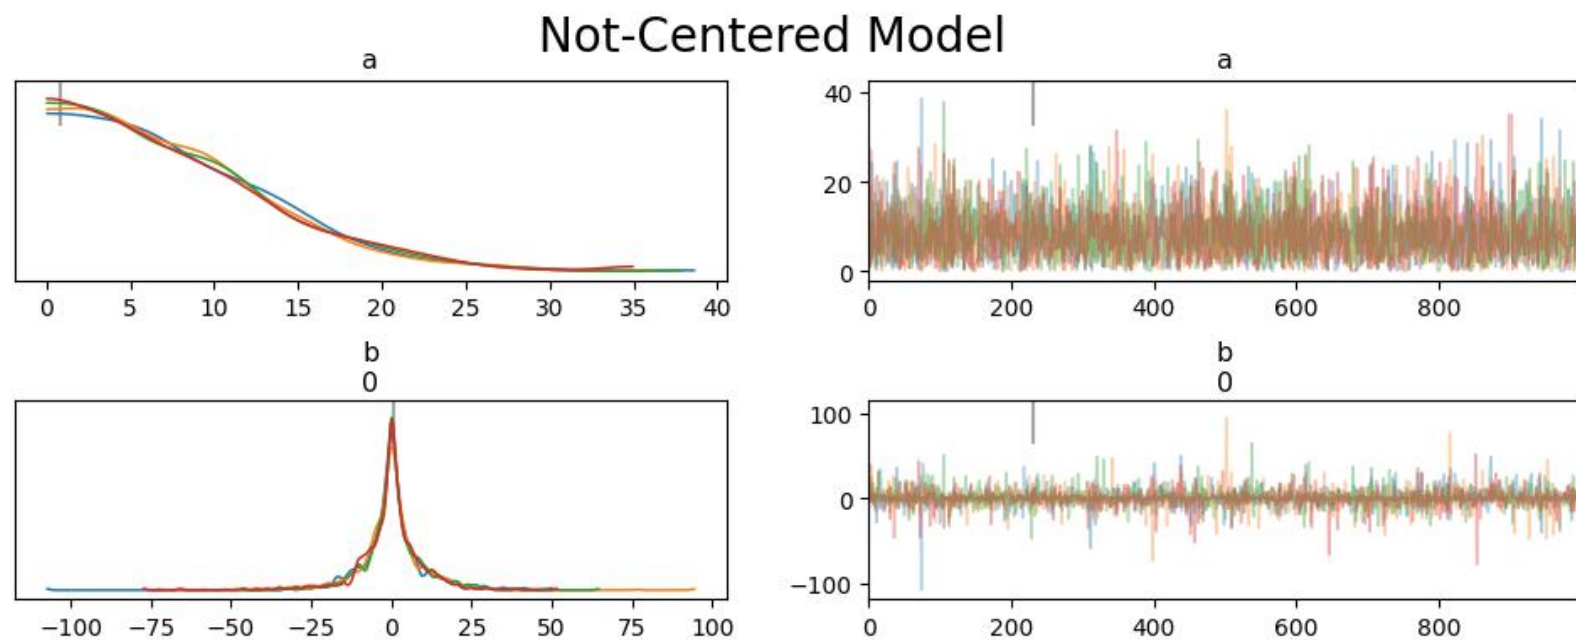
Better:



Less good:



Better:



Convergence

- How do we check convergence?
 - Look at the trace plots and see if the chains look similar
 - Overlapping KDE chains
 - Trace plots that look noisy, mixed and like they explored the different values.
 - The overlapping and mixing means that even when the different chains start from different points, they all describe the same distribution.

Convergence

- How do we check convergence?
 - Rhat (\hat{R})
 - Compare the variance between chains with the variance within the chains.
 - Ideally should be 1 (the variance between the chains is not more than the variance within the chains)
 - Values below 1.1.

Convergence

- How do we check convergence?
 - Look at the trace plots and see if the chains look similar
 - Overlapping KDE chains
 - Trace plots that look noisy, mixed and like they explored the different values.
 - Rhat
 - Below 1.1
 - Effective Sample Size

Convergence

- Effective Sample Size

- How many useful draws do we have in our sample?
- Due to autocorrelation: Number of useful draws < number of samples.
- We need an effective sample size of at least 100 per chain (400 for 4 chains).
 - We'll use at least 500
- The effective sample size is the number of samples, corrected for the autocorrelation.

$$N_{\text{eff}} = \frac{N_{\text{MCMC}}}{\sum_{\delta=-\infty}^{\infty} \rho(\delta)}$$

Convergence

- How do we check convergence?
 - Look at the trace plots and see if the chains look similar
 - Overlapping KDE chains
 - Trace plots that look noisy, mixed and like they explored the different values.
 - Rhat
 - Below 1.1
 - Effective Sample Size
 - At least 100 per chain
 - Monte Carlo Standard Error

Convergence

- Monte Carlo Standard Error

- Width of the posterior divided by the square root of the effective sample size.

$$MCSE = \frac{\sigma_{\theta}}{\sqrt{N_{\text{eff}}}}$$

- Indication of how well we are measuring the parameter.
 - Should be less than 1%-2%.

Convergence

- How do we check convergence?
 - Look at the trace plots and see if the chains look similar
 - Overlapping KDE chains
 - Trace plots that look noisy, mixed and like they explored the different values.
 - Rhat
 - Below 1.1
 - Effective Sample Size
 - At least 100 per chain
 - Monte Carlo Standard Error
 - Below 1-2%.

	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat
Less good:					
thet	0.005	0.002	1492.0	2450.0	1.0

	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat	
Better:	thet	0.004	0.002	1843.0	2899.0	1.0

Convergence

- What do we do if there is no convergence?
 - Check our codes.
 - Increase number of samples.
 - Increase number of tuning samples.
 - Change the model
 - Try a simpler model
 - Change the priors
 - Change the parameters

Divergence

- If there are divergences, it indicates that the sampler has likely found a region of high curvature in the posterior that cannot be explored properly.
 - Only for the NUTS sampler (based on the Hamiltonian Monte Carlo method).
 - We can try and reparametrize the model.
 - We can increase target accept in `pm.sample`.

Divergence

Progress	Draws	Divergences	Step size	Grad evals	Sampling Speed	Elapsed	Remaining
<div></div>	2000	13	1.23	3	2652.35 draws/s	0:00:00	0:00:00
<div></div>	2000	19	1.01	1	1383.77 draws/s	0:00:01	0:00:00
<div></div>	2000	9	1.85	1	833.34 draws/s	0:00:02	0:00:00
<div></div>	2000	7	1.73	1	574.27 draws/s	0:00:03	0:00:00

ERROR:pymc.stats.convergence:There were 48 divergences after tuning. Increase `target_accept` or reparameterize.

- That was for a default target_accept = 0.8
- Increase to 0.9

```
#adjust target accept
with pm.Model(coords=coords) as model_3:
    thet = pm.Uniform('thet', lower=-0.1, upper=1)
    y = pm.Bernoulli('y', p=thet, observed=data, dims = 'data')
    idata3 = pm.sample(1000, chains = 4, target_accept = 0.9)
```

Progress	Draws	Divergences	Step size	Grad evals	Sampling Speed	Elapsed	Remaining
<div></div>	2000	4	0.99	3	1580.85 draws/s	0:00:01	0:00:00
<div></div>	2000	3	0.77	1	832.91 draws/s	0:00:02	0:00:00
<div></div>	2000	5	1.00	3	547.14 draws/s	0:00:03	0:00:00
<div></div>	2000	3	1.40	1	438.76 draws/s	0:00:04	0:00:00

ERROR:pymc.stats.convergence:There were 15 divergences after tuning. Increase `target_accept` or reparameterize.

Divergence

```
with pm.Model(coords=coords) as model_3:  
    thet = pm.Uniform('thet', lower=-0.1, upper=1)  
    y = pm.Bernoulli('y', p=thet, observed=data, dims = 'data')  
    idata3 = pm.sample(1000, chains = 4, target_accept = 0.99)
```

Progress	Draws	Divergences	Step size	Grad evals	Sampling Speed	Elapsed	Remaining
<div></div>	2000	0	0.43	15	675.87 draws/s	0:00:02	0:00:00
<div></div>	2000	0	0.35	7	485.84 draws/s	0:00:04	0:00:00
<div></div>	2000	0	0.42	15	372.43 draws/s	0:00:05	0:00:00
<div></div>	2000	0	0.33	3	303.08 draws/s	0:00:06	0:00:00

